

CNN sensitivity analysis for land cover map models using sparse and heterogeneous satellite data

The Iberoamerican Congress on Pattern Recognition, 2024

Sebastian Moreno, Javier Lopatin, Diego Corvalán, Alejandra Bravo-Diaz.

Universidad Adolfo Ibáñez

Magister en Ciencias de la ingeniería

Doctorado en Data Science

November, 2024

Context

- Currently, it is possible to **observe** the Earth with different **remote sensing** techniques such as satellites and drones.

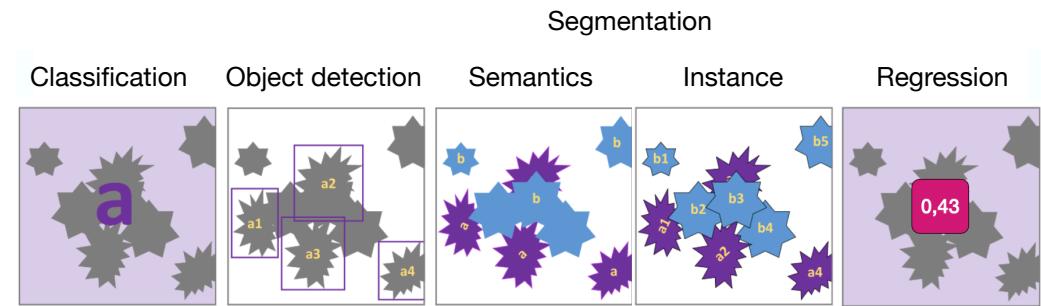
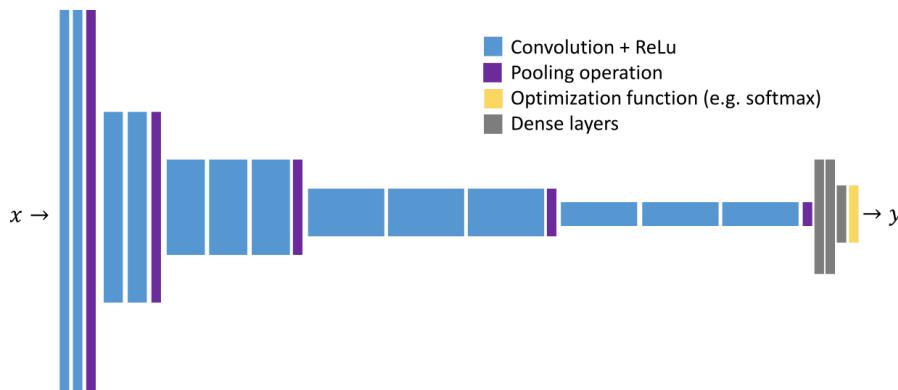


- **Land cover maps** are essential for characterizing the Earth's surface at a specific location and time.



Context

- **Convolutional neural networks** (CNNs) have transformed image processing by enabling models to learn patterns directly from images.
- **Segmentation** is one of the most common approaches for **land cover** analysis.
- **Pre-trained networks** are widely applied to address challenges in remote sensing.



Problem

Is it feasible to use pre-trained models
for land cover tasks using satellite data?

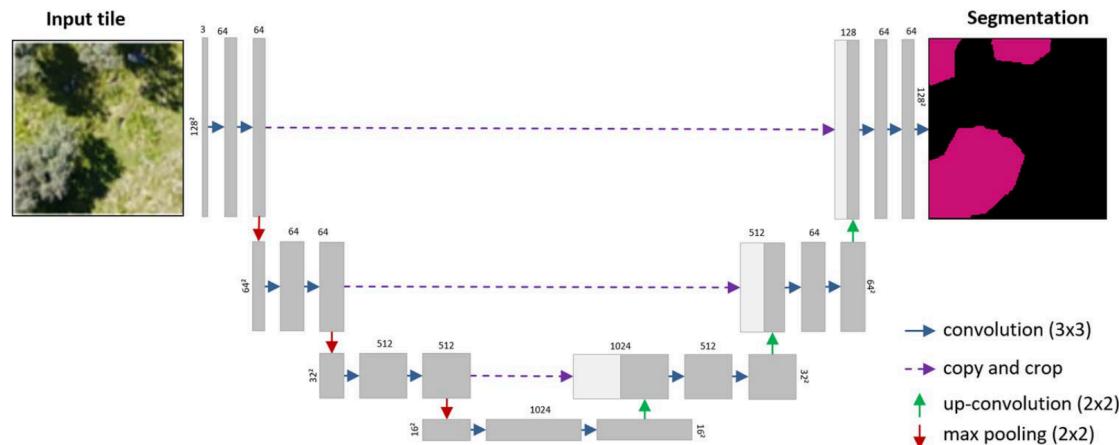
ResNet50

Xception

InceptionV3

Seresnet18

EfficientNet

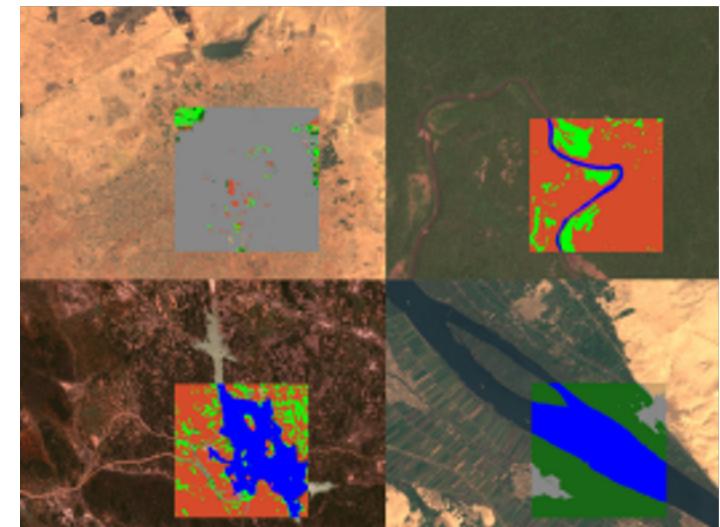


We analyze the segmentation capability of different models
and architectures on the **LandCoverNet** dataset.

Methods - Data acquisition

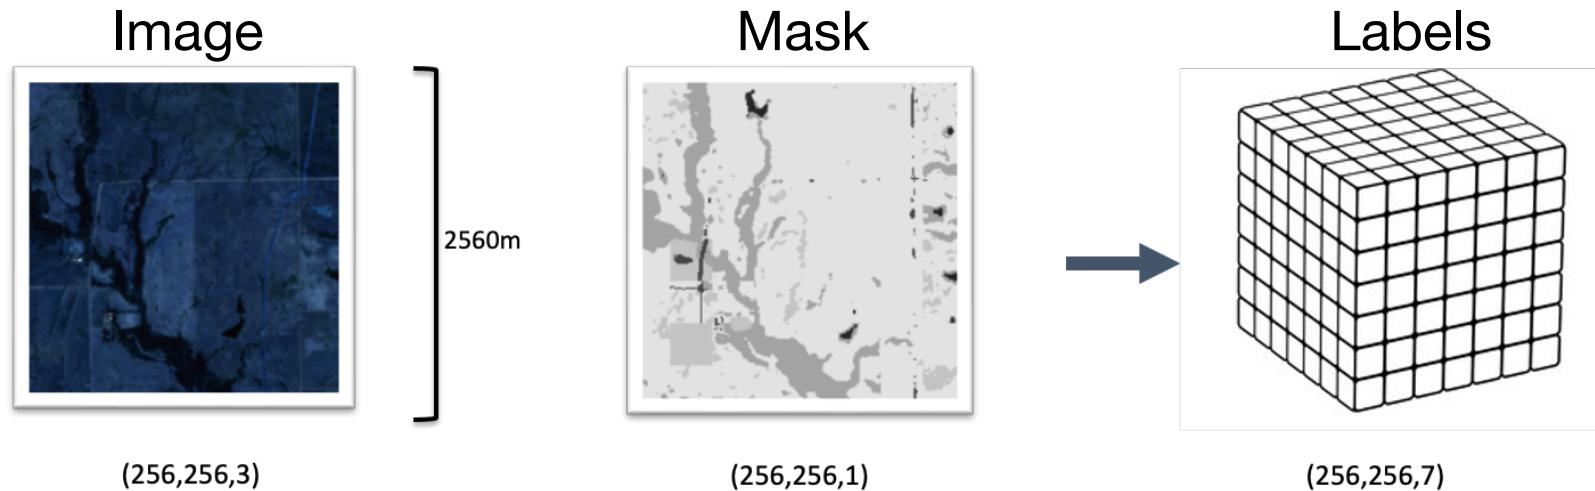
- **LandCoverNet** dataset
- The dataset was reduced to 600 images of 256x256x3, divided among Africa, South America, and North America.

| Class | % |
|-----------------------------|----------|
| • C1-Water | • 4.69% |
| • C2-Artificial bare soil | • 2.78% |
| • C3-Natural bare soil | • 10.19% |
| • C4-Snow/Ice | • 0.70% |
| • C5-Woody vegetation | • 24.07% |
| • C6- Cultivated vegetation | • 20.89% |
| • C7-Natural vegetation | • 36.68% |

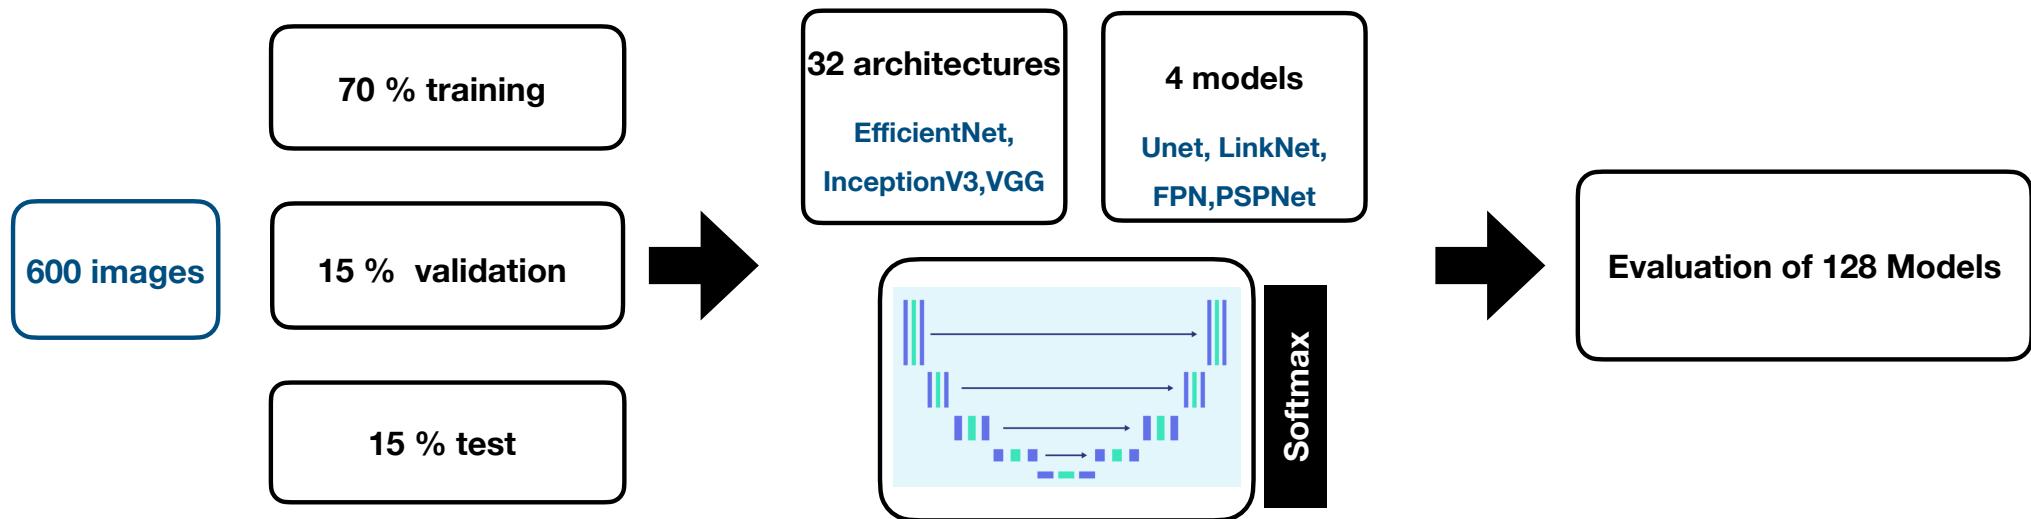


Methods - Image processing

- RGB channels were used.
- **Clouds were filtered** using the CLD channel.
- A total of **600 images** were selected (200 per continent).
- **One-hot encoding** was applied to the masks (for 7 classes).



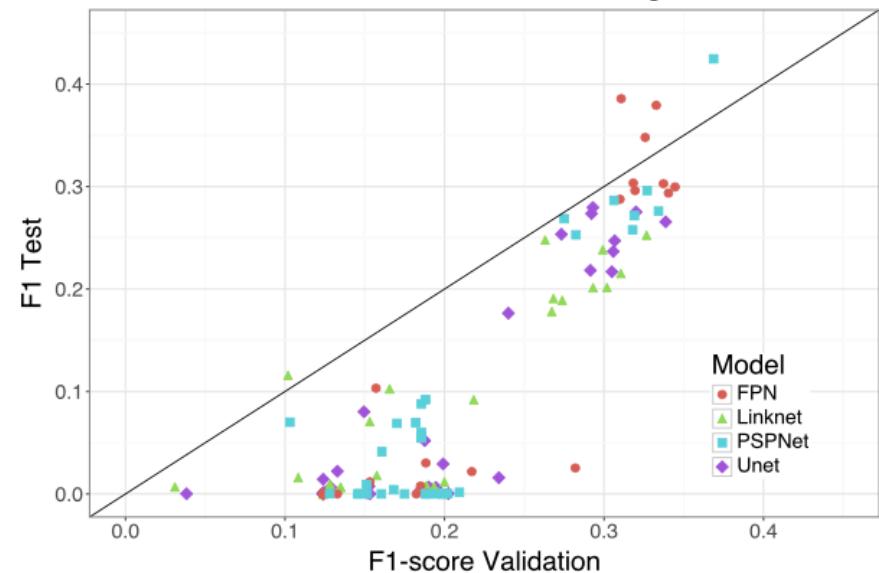
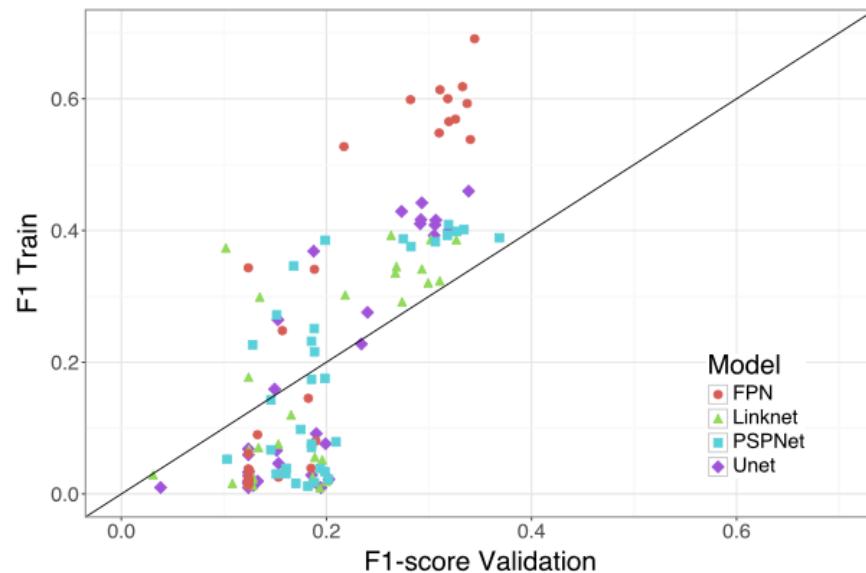
Methods - Modeling



- **128 models** were evaluated for land cover segmentation.
(4 models and 32 architectures per model).
- The **F1-score metric** was used to evaluate model performance.

Results - F1-scores

Several models from the FPN architecture exhibited overfitting, **PSPNet** and **LinkNet** showed less overfitting, with **PSPNet** performing the best.

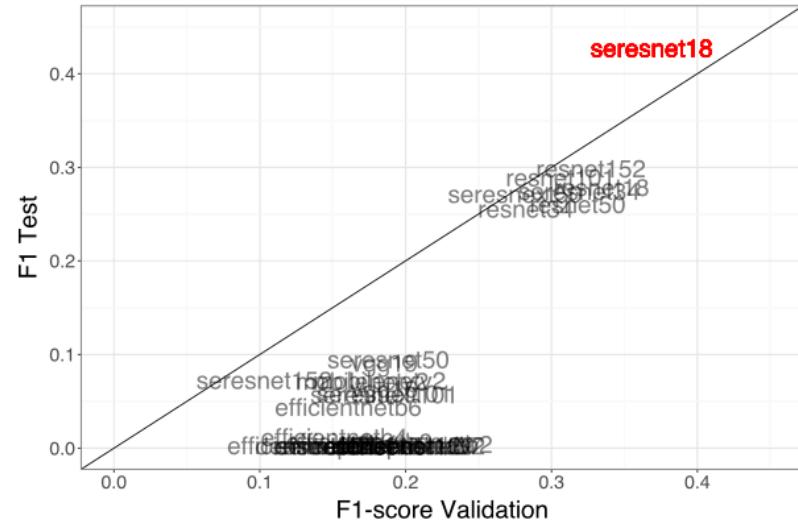
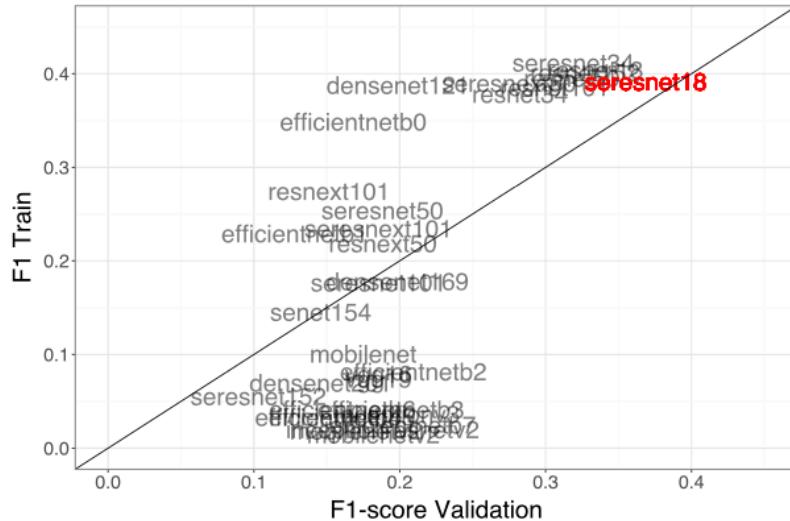


Comparing the 128 trained models.

Left: Training versus validation data. Right: Test versus Validation data.

Results - F1-scores => PSPNet model

seresnet18 has the highest F1-score in the validation data with 0.3752, and close values for training (0.3889) and test (0.4247).



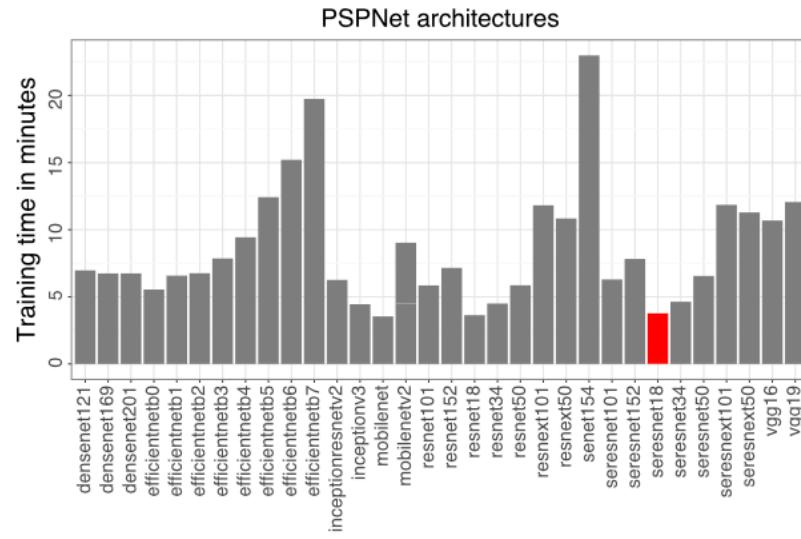
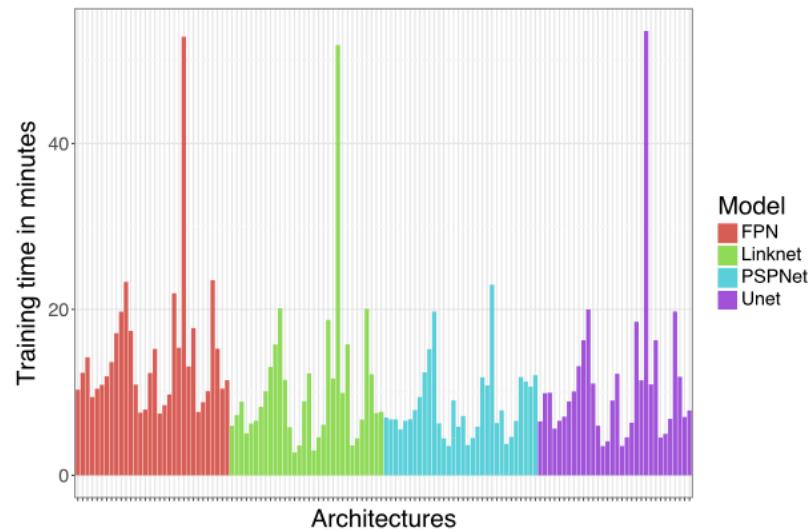
F1-scores results for the **32 architectures of the PSPNet model**.

Left: Training versus validation data. Right: Validation versus Test data.

Results - training times

Training time depends on the architectures rather than the model.

PSPNet-seresnet18 (red bar) is the fastest model and has the best F1-score.



Training times in minutes for the 128 models.

Left: Training time, we get rid of architecture names to avoid cluttering.

Right: training time of the 32 PSPNet architectures.

Results - fine tuning

The highest F1-score for validation/test was achieved with **80 unfrozen layers**.

Increasing the number of unfrozen layers could lead to overfitting.

| Unfrozen layers | Training time (mins) | F1-score | F1-score | F1-score |
|--------------------|-------------------------|---------------|---------------|---------------|
| | | Train | Validation | Test |
| 0 | 3.701 | 0.3889 | 0.3752 | 0.4247 |
| 5 | 5.506 | 0.4944 | 0.4281 | 0.4266 |
| 10 | 5.194 | 0.5187 | 0.4439 | 0.4472 |
| 20 | 5.025 | 0.5391 | 0.4332 | 0.4303 |
| 40 | 4.683 | 0.5049 | 0.4607 | 0.4551 |
| 60 | 4.326 | 0.5366 | 0.4583 | 0.4550 |
| 80 | 4.094 | 0.5253 | 0.4798 | 0.4786 |
| 100 | 2.927 | 0.4752 | 0.4537 | 0.4541 |

Fine-tuning results for the **PSPNet-seresnet18**.

The experiment unfreezes the last encoder layers and gradually increasing the number of unfrozen layers up to 100 layers.

Results - data augmentation

Data augmentation increased the training time, and **improved the performance**.

As expected, fine-tuning and data augmentation obtained the best results.

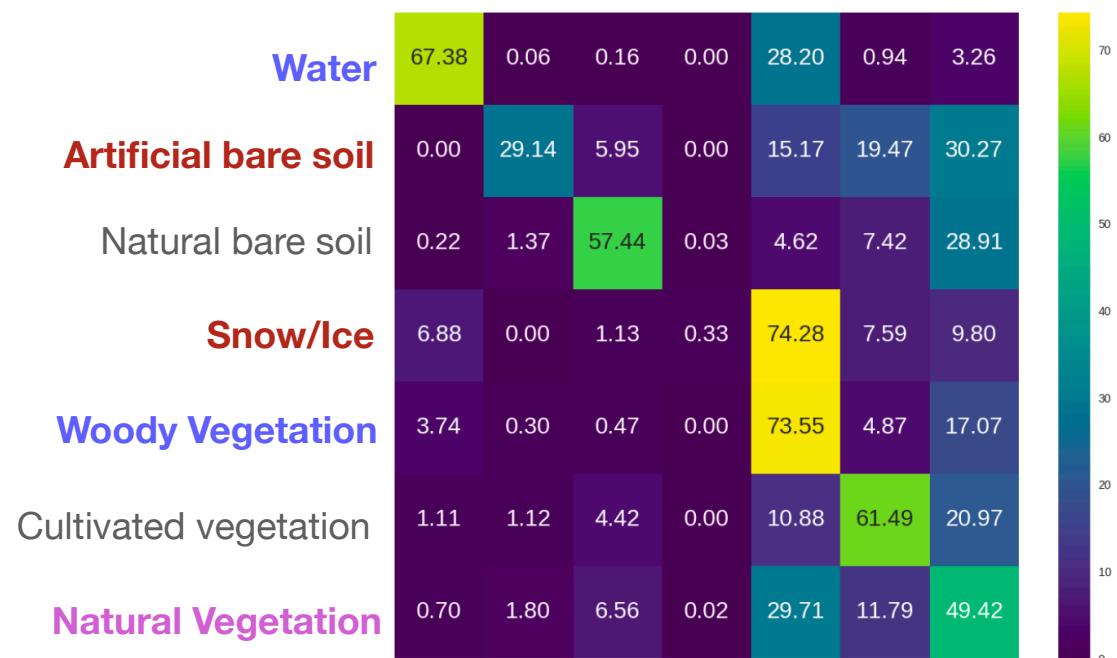
| Model | Training time (mins) | F1-score | | |
|-----------------------------------|-------------------------|---------------|---------------|---------------|
| | | Train | Validation | Test |
| Original | 3.701 | 0.3889 | 0.3650 | 0.4247 |
| Augmentation | 9.920 | 0.3728 | 0.3552 | 0.4620 |
| Fine-tuning | 4.094 | 0.5253 | 0.4798 | 0.4786 |
| Augmentation + Fine-tuning | 9.023 | 0.5314 | 0.5157 | 0.4950 |

Data augmentation and fine-tuning results for the **PSPNet-seresnet18**.

The presented models are Original, Augmentation (Original + data augmentation),
Fine-tuning (Original with fine-tuning), and Augmentation + Fine-tuning.

Results - confusion matrix

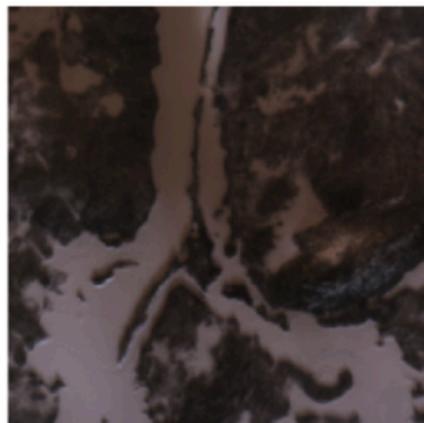
- “Water” and “Woody vegetation” have the best performances.
- Natural Vegetation is mainly confused with Woody vegetation.
- “Artificial bare soil” and “Snow/Ice” have few training data points (less than 3%) and show the worst performances.



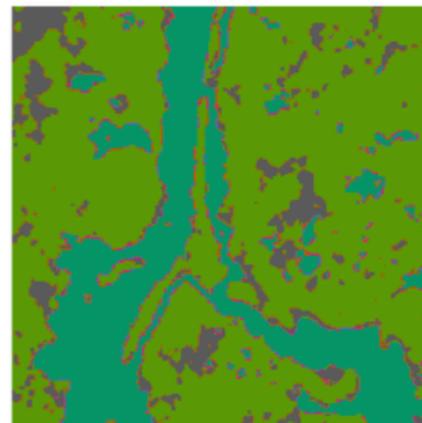
Confusion matrix for **PSPNet-seresnet18** (80 unfrozen layers and data augmentation).

Results - segmentation example

PSPNet-Seresnet18 **classifies correctly most of the water and snow/ice** (dark green and pink pixels), but **fails to classify the many small patches of natural vegetation** (gray pixels)



Original RGB image



Original mask



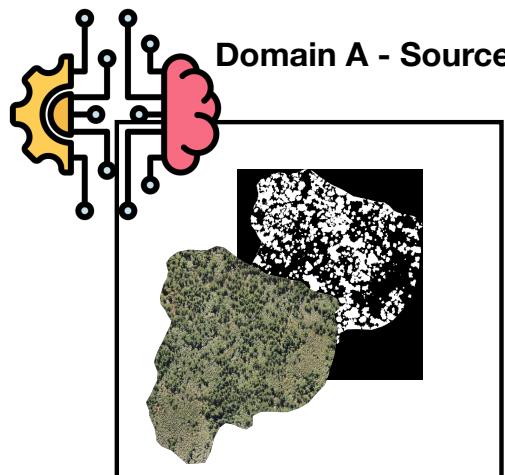
Final classification
PSPNet-Seresnet18

Conclusions

- **PSPNet-seresnet18** outperformed other models based on the F1-score.
- **Fine-tuning and data augmentation** improved the performance of the selected model (PSPNet-seresnet18).
- The confusion matrix revealed significant differences in the **performance across classes**.
- These results **are lower than other Sentinel-2-based land cover** analyses using semantic segmentation techniques.

Future work

- In remote sensing it is common to use a model learned in a specific domain and applied in another domain.
- Given the poor performance obtained in the model, can we use a model learned from another domain and applied on another dataset?



Use the model learned with data from **domain A**
to predict in **domain B**.

